Consistency, Robustness and Sparsity for Learning Algorithms

Konsistenz, Robustheit und Dünnbesetztheit von Lern-Algorithmen

Der Naturwissenschaftlichen Fakultät
der
Friedrich-Alexander-Universität Erlangen-Nürnberg

zur Erlangung des Doktorgrades
Dr. rer. nat.

vorgelegt von
Tim Roith
aus
Amberg

Als Dissertation genehmigt von der Naturwissenschaftlichen Fakultät der
Friedrich-Alexander-Universität Erlangen-Nürnberg

Tag der mündlichen Prüfung:	—
Vorsitzender des Promotionsorgans:	—
Gutachter*in:	Martin Burger
	Dejan Slepčev
	Franca Hoffmann

Acknowledgement

Coming soon...

Contents

List of Figures

Preface

This work is structured into two main parts, Part I the presentation and explanation of the topics and results presented in **??**, the peer-reviewed articles.

Grafik anpassen

Part I: Exposition	**??**: ??
Chapter 2: Learning Paradigms	—
??: ??	**????**
??: ??	**????**

Part I consists of three chapters, of which the first explains the paradigms, *unsupervised*, *semi-supervised* and *supervised* learning. The other chapters are the split up thematically, concerning the topics semi-supervised and supervised learning respectively. In each of these chapters a short introduction provides the necessary framework allowing us to explain the main contributions. The following publications are reprinted in **??**:

[LIP-I] T. Roith and L. Bungert. "Continuum limit of Lipschitz learning on graphs." In: *Foundations of Computational Mathematics* (2022), pp. 1–39.

[LIP-II] L. Bungert, J. Calder, and T. Roith. "Uniform convergence rates for Lipschitz learning on graphs." In: *IMA Journal of Numerical Analysis* (Sept. 2022). DOI: 10.1093/imanum/drac048.

[CLIP] L. Bungert, R. Raab, T. Roith, L. Schwinn, and D. Tenbrinck. "CLIP: Cheap Lipschitz training of neural networks." In: *Scale Space and Variational Methods in Computer Vision: 8th International Conference, SSVM 2021, Virtual Event, May 16–20, 2021, Proceedings*. Springer. 2021, pp. 307–319.

[BREG-I] L. Bungert, T. Roith, D. Tenbrinck, and M. Burger. "A Bregman learning framework for sparse neural networks." In: *Journal of Machine Learning Research* 23.192 (2022), pp. 1–43.

[FNO] S. Kabri, T. Roith, D. Tenbrinck, and M. Burger. "Resolution-Invariant Image Classification based on Fourier Neural Operators." In: *Scale Space and Variational Methods in Computer Vision: 9th International Conference, SSVM 2023, Proceedings*. Springer. 2023, pp. 307–319.

The following two works that are not part of this thesis but provide an additional insight.

[LIP-III] L. Bungert, J. Calder, and T. Roith. *Ratio convergence rates for Euclidean first-passage percolation: Applications to the graph infinity Laplacian*. 2022. arXiv: 2210.09023 [math.PR].

[BREG-II] L. Bungert, T. Roith, D. Tenbrinck, and M. Burger. "Neural Architecture Search via Bregman Iterations." In: (2021). arXiv: 2106.02479 [cs.LG].

TR's Contribution

Here we list TR's contribution to the publications included in the thesis.

[LIP-I]: This work builds upon the findings in TR's master thesis [Roi21]. It is however important to note that the results constitute a significant extension and are conceptually stronger than the ones in [Roi21], see **??**. TR adapted the continuum limit framework to the L^∞ case, worked out most of the proofs and wrote a significant part of the paper. In collaboration with LB, he identified the crucial domain assumptions that allow to work on non-convex domains and proved convergence for approximate boundary conditions.

[LIP-II]: In collaboration with LB, TR worked on the convergence proofs building upon the ideas of JC. He contributed to both the numeric and the analysis conducted in the paper.

[CLIP]: TR worked out the main algorithm proposed in the paper together with LB, based on LB's idea. Together with LS and RR he conducted the numerical examples and also wrote most of the source code. Furthermore, he wrote large parts of the paper.

[BREG-I]: TR expanded LB's ideas of employing Bregman iteration for sparse training. Together with MB and LB he worked out the convergence analysis of stochastic Bregman iterations. Here, he also proposed a profound sparse initialization strategy. Furthermore, he conducted the numerical examples and wrote most of the source code.

[FNO]: This work is based on SK's masters thesis, employing the initial ideas of MB for resolution invariance with FNOs. In the paper TR worked out the proofs for well-definedness and Fréchet-differentiability, together with SK. He wrote large parts of the paper and the source code. Here, he conducted the numerical studies in collaboration with SK.

Part I.

Exposition

Chapter 1
Introduction

The field of *machine learning* emerged in the 1950s [Sam59; Ros58], motivated by the idea of letting a machine discover algorithms and patterns without having to explicitly arrange them by hand. After the initial phase and multiple "AI-winters" [SG96], numerous important developments—e.g. the rediscovery of the backpropagation algorithm, originally due to [Kel60; Ros+62] and then popularized in [RHW86], see e.g. [Sch22]—contributed to the relevance of learning methods. The advances in computer hardware together with the availability of large amounts of data, finally allowed the machine learning enthusiasm of the recent years to spark. While "deep" learning methods—i.e. techniques involving many stacked neural layers as originally proposed in [Ros58]—are the most prominent examples, there is a whole zoo of learning-based strategies that are actively applied in fields like computer vision [Cha+21], natural language processing [Khu+23] or healthcare [She+22]. In this work we mainly focus on data-driven approaches, applied to classification tasks, where the concrete modality of the given data determines our approach. Namely, we focus on supervised—the dataset consists of input-output pairs, i.e. is fully labeled—and the semi-supervised—the data is only partially labeled—learning tasks.

For both regimes especially the last 20 years have seen great success of these data-driven methods. However, the sometimes purely heuristic learning strategies also exhibit serious drawbacks. In the supervised setting one is usually interested in the generalization behavior of a learned classifier, i.e. how good is the performance on unseen inputs that are not part of the given training data. Unfortunately in [GSS14] it was discovered, that this performance can be completely corrupted, by small, seemingly invisible perturbations known as *adversarial attacks*. More generally this phenomenon leads us to the issue of *input robustness*. Given some input x and suppose that a human and some machine would classify this input to be of type c. In a rather vague but demonstrative formulation the key implication for transformed input $\bar{x}$ we want to obtain is

$$\left.\begin{array}{r}\bar{x} \text{ is close to } x,\\ \bar{x} \text{ is still classified as } c \text{ by a human}\end{array}\right\} \Rightarrow \text{the machine classifies } \bar{x} \text{ as } c.$$

Next to adversarial examples this also includes resolution changes of images, which do not—if they are reasonably small—change the classification by a human. For the semi-supervised setting

weiter schreiben

Chapter 2
Learning Paradigms

Throughout this thesis, we assume to be given data $\mathcal{X}_n \subset \mathcal{X} \subset \mathbb{R}^d$ consisting of n data points. We consider task of *learning* a function $f : \tilde{\mathcal{X}} \to \mathcal{Y}$ from the given data, where the two most important cases for us are

- **classification**: f assigns a label to each $x \in \tilde{\mathcal{X}}$ out of a total of $C \in \mathbb{N}$ possible classes, i.e. $\mathcal{Y} = \{1, \dots, C\}$. In some architectures the last layer of the neural network is given as a vector $y \in \mathbb{R}^C$. Typically, this vector is a probability vector, i.e.

$$y \in \Delta^C := \left\{ z \in [0,1]^d : \sum_{i=1}^{C} z_i = 1 \right\}.$$

  This can be enforced via the softmax function [Bri90] $\text{softmax} : \mathbb{R}^d \to \mathbb{R}^d$

$$\text{softmax}(z)_i := \frac{\exp(z_i)}{\sum_{j=1}^{C} \exp(z_j)}$$

  which was actually introduced by Boltzman in [Bol68]. This allows the interpretation that the ith entry of $f_\theta(x) \in \Delta^C$ models the probability that x belongs to class i. In order to obtain a label one can simply choose the maximum entry, $\text{argmax}_{i=1,\dots,C} f_\theta(x)_i$.

- **image denoising**: f outputs a denoised version of an input image. Here we have $\mathcal{X} = \mathcal{Y} = \mathbb{R}^{K \times N \times M}$, where
  - $K \in \mathbb{N}$ is the number of color channels,
  - N, M denote the width and height of the image.

The set $\tilde{\mathcal{X}} \subset \mathbb{R}^d$ is usually either the set of data points $\mathcal{X}_n$ or the whole space $\mathcal{X}$. The learning paradigms we consider in this thesis, differ by their usage of labeled data. We review the concepts in the following.

2.1. Unsupervised Learning

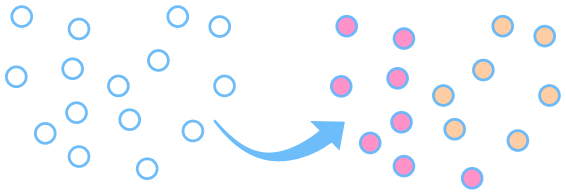

In this case we are not given any labeled data. In our context the most important application is data clustering. Other tasks involve dimensionality reduction or density estimation, see [ST14]. The clustering task consists of grouping data based on some similarity criterion. In this sense, clustering can also be interpreted as classification, i.e., the desired function is a mapping $f : \tilde{\mathcal{X}} \to \{1, \dots, C\}$ where $C \in \mathbb{N}$ denotes the number of clusters. Typically, one wants to obtain a clustering of the given data set, i.e., $\tilde{\mathcal{X}} = \mathcal{X}_n$. We list some of the typical clustering methods below:

- K-means algorithm [Ste+56],
- Expectation Maximization [DLR77],
- Cheeger cuts [GS15; SB09; Gar+16; GMT22],
- spectral clustering [GS18; THH21; Hof+22].

Unsupervised learning is not the main focus of this present work. However, we note that especially the concepts developed in [GS15] for Cheeger cuts are crucial for the continuum limit framework in **??**.

2.2. Supervised Learning

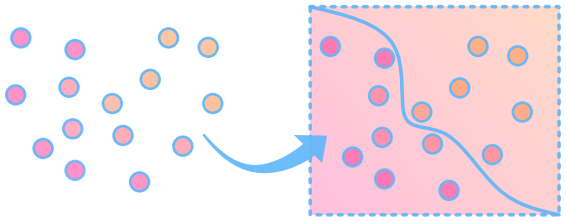

In this setting, each data point $x \in \mathcal{X}_n$ is labeled, via a given function $g : \mathcal{X}_n \to \mathcal{Y}$ such that we have a finite training set $\mathcal{T} = \{(x, g(x)) : x \in \mathcal{X}_n\}$. The task is then to infer a function defined on the underlying space, i.e. $f : \mathcal{X} \to \mathcal{Y}$, i.e. we want to assign a label to unseen $x \in \mathcal{X}$ that are not necessarily part of the given data. Often, one models the problem via a joint probability function $P_{\mathcal{X},\mathcal{Y}}$ and assumes that the training data are i.i.d. w.r.t. $P_{\mathcal{X},\mathcal{Y}}$. In this interpretation, a neural network can aim to model the conditional $P(y|x)$ for an input $x \in \mathcal{X}$ and output $y \in \mathcal{Y}$.

In order to *learn* the function f from the given data, one needs to choose a parameterized class of functions $\mathcal{U}$, where typically each element can be describe by a finite number of parameters. Among others, common methods or parametrizations include

- Support vector machines [CV95; SS05],
- decision Trees [MS63; Bre+84],
- neural networks [Tur04; Ros58; MP69].

In ?? we exclusively focus on supervised learning algorithms employing neural networks. We refer to [Sch15] for an exhaustive historical overview. The concrete setting and learning framework is given in ??.

2.3. Semi-Supervised Learning

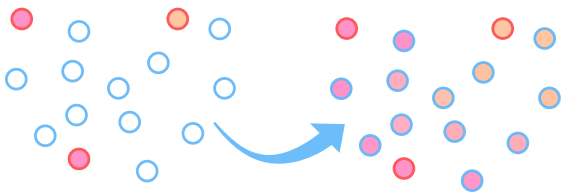

In the semi-supervised setting we assume that only a fraction of the data $\mathcal{X}_n$ is labeled, i.e., we are given a a function $g : \mathcal{O}_n \to \mathcal{Y}$ where $\mathcal{O}_n \subset \mathcal{X}_n$ is the set of labeled data. Typically the labeled data constitutes only a small fraction of all available points, i.e. $|\mathcal{O}_n| << |\mathcal{X}_n|$. In this thesis we restrict ourselves to the *transductive setting*, i.e. we want to infer a function acting only on the data $f : \mathcal{X}_n \to \mathcal{Y}$. This is opposed to the inductive setting, where f also classifies unseen points $x \in \mathcal{X}$, [Zhu05]. Common algorithms and methods include

- expectation maximization and mixture models [DLR77; CCC+03],
- self-training and co-training [BM98],
- graph-based learning [Zhu05].

Mostly, we consider the extension task with $\mathcal{Y}$ being chosen as $\mathbb{R}$. In application this can be seen as a binary classification task, where for $o \in \mathcal{O}_n$ we have $g(o) = 1$ if o belongs to a some class and $g(o) = 0$ otherwise. The function $f : \mathcal{X}_n \to \mathbb{R}$ then determines the probability that any vertex $x \in X_n$ belongs o this class, where we can binarize the output afterwards via some thresholding, e.g.,

$$x \text{ belongs to the class } \Leftrightarrow f(x) > 0.5$$

This methodology can be extended to classification tasks beyond the binary case, via the so-called one-vs-all technique [ZGL03]. Given a classification problem with $C \in \mathbb{N}$ possible classes, we assume that the labeling function $g : \mathcal{O}_n \to \Delta^C$ outputs one-hot vectors, i.e. $g(o)_c = 1$ if o belongs to class c and $g(o)_c = 0$ otherwise, for every $c = 1, \dots, C$. We then perform the binary classification problem "x belongs to class c" for every $c = 1, \dots, C$, by considering the extension task of

$$g_c : \mathcal{O}_n \to \mathbb{R} \qquad g_c(o) = g(o)_c,$$

which yields a function $f_c(o)$. The final output can then either obtained by taking the argmax, i.e. $f : \mathcal{X}_n \to \{1, \dots, C\}$

$$f(x) := \operatorname{argmax}_{c=1,\dots,C} f_c(x)$$

or by applying a softmax to obtain a probability vector, i.e. $f : \mathcal{X}_n \to \Delta^C$

$$f(x) := \operatorname{softmax}\left(f_1(x), \dots, f_c(x)\right).$$

In ?? we focus on graph-based learning algorithms, however we refer to [Zhu05] for a overview of semi-supervised learning algorithms.

Bibliography

Books

[Ros+62] F. Rosenblatt et al. *Principles of neurodynamics: Perceptrons and the theory of brain mechanisms*. Vol. 55. Spartan books Washington, DC, 1962.

[Bre+84] L. Breiman, J. Friedman, R. Olshen, and C. Stone. *Clasificaion and regression trees*. CRC Press, New York, 1984.

[Zhu05] X. Zhu. *Semi-supervised learning with graphs*. Carnegie Mellon University, 2005.

[Lin17] P. Lindqvist. *Notes on the p-Laplace equation*. 161. University of Jyväskylä, 2017.

[AF03] R. A. Adams and J. J. Fournier. *Sobolev spaces*. Elsevier, 2003.

[Eva18] L. Evans. *Measure theory and fine properties of functions*. Routledge, 2018.

[BB11] H. Brezis and H. Brézis. *Functional analysis, Sobolev spaces and partial differential equations*. Vol. 2. 3. Springer, 2011.

[Sch69] J. T. Schwartz. *Nonlinear Functional Analysis* -. Boca Raton, Fla: CRC Press, 1969.

[Lin16] P. Lindqvist. *Notes on the infinity Laplace equation*. Springer, 2016.

[Bra02] A. Braides. *Gamma-convergence for Beginners*. Vol. 22. Oxford University Press, Oxford, 2002.

[Dal12] G. Dal Maso. *An introduction to Γ-convergence*. Vol. 8. Springer Science & Business Media, 2012.

[Hau14] F. Hausdorff. *Grundzüge der Mengenlehre*. Viet, Leipzig, 1914.

[DS88] N. Dunford and J. T. Schwartz. *Linear operators, part 1: general theory*. Vol. 10. John Wiley & Sons, 1988.

[COR98] G. Cybenko, D. P. O'Leary, and J. Rissanen. *The mathematics of information coding, extraction and distribution*. Vol. 107. Springer Science & Business Media, 1998.

[Dac07] B. Dacorogna. *Direct methods in the calculus of variations*. Vol. 78. Springer Science & Business Media, 2007.

[SB14] S. Shalev-Shwartz and S. Ben-David. *Understanding machine learning: From theory to algorithms*. Cambridge university press, 2014.

[VD95] G. Van Rossum and F. L. Drake Jr. *Python reference manual*. Centrum voor Wiskunde en Informatica Amsterdam, 1995.

[Roc97] R. Rockafellar. *Convex analysis*. Princeton, N.J: Princeton University Press, 1997.

[BC11] H. Bauschke and P. Combettes. *Convex analysis and monotone operator theory in Hilbert spaces*. New York: Springer, 2011.

[Eul24] L. Euler. *Institutionum calculi integralis*. Vol. 1. impensis Academiae imperialis scientiarum, 1824.

[BV04] S. P. Boyd and L. Vandenberghe. *Convex optimization*. Cambridge university press, 2004.

[GBC16] I. Goodfellow, Y. Bengio, and A. Courville. *Deep Learning*. http://www.deeplearningbook.org. MIT Press, 2016.

[Ral81] L. B. Rall. *Automatic differentiation: Techniques and applications*. Springer, 1981.

[GW87] R. C. Gonzales and P. Wintz. *Digital image processing*. Addison-Wesley Longman Publishing Co., Inc., 1987.

[Trö10] F. Tröltzsch. *Optimal Control of Partial Differential Equations: Theory, Methods, and Applications*. Vol. 112. Graduate Studies in Mathematics. American Mathematical Society, Providence, Rhode Island, 2010.

[Gra14] L. Grafakos. *Classical Fourier Analysis*. 3rd ed. Graduate Texts in Mathematics. Springer, New York, NY, 2014.

[AP93] A. Ambrosetti and G. Prodi. *A Primer of Nonlinear Analysis*. Cambridge University Press, 1993.

Articles

[LIP-I] T. Roith and L. Bungert. "Continuum limit of Lipschitz learning on graphs." In: *Foundations of Computational Mathematics* (2022), pp. 1–39.

[LIP-II] L. Bungert, J. Calder, and T. Roith. "Uniform convergence rates for Lipschitz learning on graphs." In: *IMA Journal of Numerical Analysis* (Sept. 2022). DOI: 10.1093/imanum/drac048.

[BREG-I] L. Bungert, T. Roith, D. Tenbrinck, and M. Burger. "A Bregman learning framework for sparse neural networks." In: *Journal of Machine Learning Research* 23.192 (2022), pp. 1–43.

[Sam59] A. L. Samuel. "Some studies in machine learning using the game of checkers." In: *IBM Journal of research and development* 3.3 (1959), pp. 210–229.

[Ros58] F. Rosenblatt. "The perceptron: a probabilistic model for information storage and organization in the brain." In: *Psychological review* 65.6 (1958), p. 386.

[Kel60] H. J. Kelley. "Gradient theory of optimal flight paths." In: *Ars Journal* 30.10 (1960), pp. 947–954.

[RHW86] D. E. Rumelhart, G. E. Hinton, and R. J. Williams. "Learning representations by back-propagating errors." In: *nature* 323.6088 (1986), pp. 533–536.

[Sch22] J. Schmidhuber. "Annotated history of modern AI and Deep learning." In: *arXiv preprint arXiv:2212.11279* (2022).

[Cha+21] J. Chai, H. Zeng, A. Li, and E. W. Ngai. "Deep learning in computer vision: A critical review of emerging techniques and application scenarios." In: *Machine Learning with Applications* 6 (2021), p. 100134.

[Khu+23] D. Khurana, A. Koli, K. Khatter, and S. Singh. "Natural language processing: State of the art, current trends and challenges." In: *Multimedia tools and applications* 82.3 (2023), pp. 3713–3744.

[She+22] M. Shehab, L. Abualigah, Q. Shambour, M. A. Abu-Hashem, M. K. Y. Shambour, A. I. Alsalibi, and A. H. Gandomi. "Machine learning in medical applications: A review of state-of-the-art methods." In: *Computers in Biology and Medicine* 145 (2022), p. 105458.

[GSS14] I. J. Goodfellow, J. Shlens, and C. Szegedy. "Explaining and harnessing adversarial examples." In: *arXiv preprint arXiv:1412.6572* (2014).

[Bol68] L. Boltzmann. "Studien über das Gleichgewicht der lebenden Kraft." In: *Wissenschafiliche Abhandlungen* 1 (1868), pp. 49–96.

[ST14] A. Subramanya and P. P. Talukdar. "Graph-based semi-supervised learning." In: *Synthesis Lectures on Artificial Intelligence and Machine Learning* 8.4 (2014), pp. 1–125.

[Ste+56] H. Steinhaus et al. "Sur la division des corps matériels en parties." In: *Bull. Acad. Polon. Sci* 1.804 (1956), p. 801.

[DLR77] A. P. Dempster, N. M. Laird, and D. B. Rubin. "Maximum likelihood from incomplete data via the EM algorithm." In: *Journal of the royal statistical society: series B (methodological)* 39.1 (1977), pp. 1–22.

[GS15] N. García Trillos and D. Slepčev. "Continuum Limit of Total Variation on Point Clouds." In: *Archive for Rational Mechanics and Analysis* 220.1 (2015), pp. 193–241. DOI: 10.1007/s00205-015-0929-z.

[SB09] A. Szlam and X. Bresson. "A total variation-based graph clustering algorithm for cheeger ratio cuts." In: *UCLA Cam report* (2009), pp. 09–68.

[Gar+16] N. García Trillos, D. Slepčev, J. Von Brecht, T. Laurent, and X. Bresson. "Consistency of Cheeger and ratio graph cuts." In: *The Journal of Machine Learning Research* 17.1 (2016), pp. 6268–6313.

[GMT22] N. García Trillos, R. Murray, and M. Thorpe. "From graph cuts to isoperimetric inequalities: Convergence rates of Cheeger cuts on data clouds." In: *Archive for Rational Mechanics and Analysis* 244.3 (2022), pp. 541–598.

[GS18] N. García Trillos and D. Slepčev. "A variational approach to the consistency of spectral clustering." In: *Applied and Computational Harmonic Analysis* 45.2 (2018), pp. 239–281.

[THH21] N. G. Trillos, F. Hoffmann, and B. Hosseini. "Geometric structure of graph Laplacian embeddings." In: *The Journal of Machine Learning Research* 22.1 (2021), pp. 2934–2988.

[Hof+22] F. Hoffmann, B. Hosseini, A. A. Oberai, and A. M. Stuart. "Spectral analysis of weighted Laplacians arising in data clustering." In: *Applied and Computational Harmonic Analysis* 56 (2022), pp. 189–249.

[CV95] C. Cortes and V. Vapnik. "Support-vector networks." In: *Machine learning* 20 (1995), pp. 273–297.

[MS63] J. N. Morgan and J. A. Sonquist. "Problems in the analysis of survey data, and a proposal." In: *Journal of the American statistical association* 58.302 (1963), pp. 415–434.

[MP69] M. Minsky and S. Papert. "An introduction to computational geometry." In: *Cambridge tiass., HIT* 479 (1969), p. 480.

[Sch15] J. Schmidhuber. "Deep learning in neural networks: An overview." In: *Neural Networks* 61 (2015), pp. 85–117. DOI: https://doi.org/10.1016/j.neunet.2014.09.003.

[ST19] D. Slepcev and M. Thorpe. "Analysis of p-Laplacian regularization in semisupervised learning." In: *SIAM Journal on Mathematical Analysis* 51.3 (2019), pp. 2085–2120.

[Cal19] J. Calder. "Consistency of Lipschitz learning with infinite unlabeled data and finite labeled data." In: *SIAM Journal on Mathematics of Data Science* 1.4 (2019), pp. 780–812.

[FCL19] M. Flores, J. Calder, and G. Lerman. "Algorithms for Lp-based semi-supervised learning on graphs." In: *arXiv preprint arXiv:1901.05031* (2019).

[CT22] J. Calder and N. G. Trillos. "Improved spectral convergence rates for graph Laplacians on ε-graphs and k-NN graphs." In: *Applied and Computational Harmonic Analysis* 60 (2022), pp. 123–175.

[ACJ04] G. Aronsson, M. Crandall, and P. Juutinen. "A tour of the theory of absolutely minimizing functions." In: *Bulletin of the American mathematical society* 41.4 (2004), pp. 439–505.

[ETT15] A. Elmoataz, M. Toutain, and D. Tenbrinck. "On the p-Laplacian and ∞-Laplacian on graphs with applications in image and data processing." In: *SIAM Journal on Imaging Sciences* 8.4 (2015), pp. 2412–2451.

[NSZ09] B. Nadler, N. Srebro, and X. Zhou. "Statistical analysis of semi-supervised learning: The limit of infinite unlabelled data." In: *Advances in neural information processing systems* 22 (2009).

[AL11] M. Alamgir and U. Luxburg. "Phase transition in the family of p-resistances." In: *Advances in neural information processing systems* 24 (2011).

[VBB08] U. Von Luxburg, M. Belkin, and O. Bousquet. "Consistency of spectral clustering." In: *The Annals of Statistics* (2008), pp. 555–586.

[GK06] E. Giné and V. Koltchinskii. "Empirical graph Laplacian approximation of Laplace-Beltrami operators: large sample results." In: *Lecture Notes-Monograph Series* (2006), pp. 238–259.

[Hof+20] F. Hoffmann, B. Hosseini, Z. Ren, and A. M. Stuart. "Consistency of semi-supervised learning algorithms on graphs: Probit and one-hot methods." In: *The Journal of Machine Learning Research* 21.1 (2020), pp. 7549–7603.

[Dun+20] M. M. Dunlop, D. Slepčev, A. M. Stuart, and M. Thorpe. "Large data and zero noise limits of graph-based semi-supervised learning algorithms." In: *Applied and Computational Harmonic Analysis* 49.2 (2020), pp. 655–697.

[CS20] J. Calder and D. Slepčev. "Properly-weighted graph Laplacian for semi-supervised learning." In: *Applied mathematics & optimization* 82 (2020), pp. 1111–1159.

[vLB04] U. von Luxburg and O. Bousquet. "Distance-Based Classification with Lipschitz Functions." In: *J. Mach. Learn. Res.* 5.Jun (2004), pp. 669–695.

[Jen93] R. Jensen. "Uniqueness of Lipschitz extensions: minimizing the sup norm of the gradient." In: *Archive for Rational Mechanics and Analysis* 123 (1993), pp. 51–74.

[Kir34] M. Kirszbraun. "Über die zusammenziehende und Lipschitzsche Transformationen." In: *Fundamenta Mathematicae* 22 (1934), pp. 77–108. DOI: 10.4064/fm-22-1-77-108.

[Whi92] H. Whitney. "Analytic extensions of differentiable functions defined in closed sets." In: *Hassler Whitney Collected Papers* (1992), pp. 228–254.

[McS34] E. J. McShane. "Extension of range of functions." In: (1934).

[Aro67] G. Aronsson. "Extension of functions satisfying Lipschitz conditions." In: *Arkiv för Matematik* 6.6 (1967), pp. 551–561.

[Per23] O. Perron. "Eine neue Behandlung der ersten Randwertaufgabe für Δu= 0." In: *Mathematische Zeitschrift* 18 (1923), pp. 42–54.

[Aro68] G. Aronsson. "On the partial differential equation $u_x^2 u_{xx} + 2u_x u_y u_{xy} + u_y^2 u_{yy} = 0$." In: *Arkiv för matematik* 7 (1968), pp. 395–425.

[Yu06] Y. Yu. "A remark on C2 infinity-harmonic functions." In: (2006).

[BDM89] T. Bhattacharya, E. DiBenedetto, and J. Manfredi. "Limits as $p \to \infty$ of $Delta_p u_p = f$ and related extremal problems." In: *Rend. Sem. Mat. Univ. Politec. Torino* 47 (1989), pp. 15–68.

[Bur48] J. M. Burgers. "A mathematical model illustrating the theory of turbulence." In: *Advances in applied mechanics* 1 (1948), pp. 171–199.

[ASS11] S. Armstrong, C. Smart, and S. Somersille. "An infinity Laplace equation with gradient term and mixed boundary conditions." In: *Proceedings of the American Mathematical Society* 139.5 (2011), pp. 1763–1776.

[AS10] S. N. Armstrong and C. K. Smart. "An easy proof of Jensen's theorem on the uniqueness of infinity harmonic functions." In: *Calculus of Variations and Partial Differential Equations* 37 (2010), pp. 381–384.

[JS06] P. Juutinen and N. Shanmugalingam. "Equivalence of AMLE, strong AMLE, and comparison with cones in metric measure spaces." In: *Mathematische Nachrichten* 279.9-10 (2006), pp. 1083–1098.

[Per+09] Y. Peres, O. Schramm, S. Sheffield, and D. Wilson. "Tug-of-war and the infinity Laplacian." In: *Journal of the American Mathematical Society* 22.1 (2009), pp. 167–210.

[NS12] A. Naor and S. Sheffield. "Absolutely minimal Lipschitz extension of tree-valued mappings." In: *Mathematische Annalen* 354 (2012), pp. 1049–1078.

[Kur22] C. Kuratowski. "Sur l'opération A de l'analysis situs." In: *Fundamenta Mathematicae* 3.1 (1922), pp. 182–199.

[ČFK66] E. Čech, Z. Frolík, and M. Katětov. "Topological spaces." In: (1966).

[DF75] E. De Giorgi and T. Franzoni. "Su un tipo di convergenza variazionale." In: *Atti della Accademia Nazionale dei Lincei. Classe di Scienze Fisiche, Matematiche e Naturali. Rendiconti* 58.6 (1975), pp. 842–850.

[Mod77] L. Modica. "Un esempio di Γ-convergenza." In: *Boll. Un. Mat. Ital. B* 14 (1977), pp. 285–299.

[CGL10] A. Chambolle, A. Giacomini, and L. Lussardi. "Continuous limits of discrete perimeters." In: *ESAIM: Mathematical Modelling and Numerical Analysis* 44.2 (2010), pp. 207–230.

[BY12] A. Braides and N. K. Yip. "A quantitative description of mesh dependence for the discretization of singularly perturbed nonconvex problems." In: *SIAM Journal on Numerical Analysis* 50.4 (2012), pp. 1883–1898.

[VB+12] Y. Van Gennip, A. L. Bertozzi, et al. "Γ-convergence of graph Ginzburg-Landau functionals." In: *Adv. Differential Equations* 17.11-12 (2012), pp. 1115–1180.

[Fré06] M. M. Fréchet. "Sur quelques points du calcul fonctionnel." In: *Rendiconti del Circolo Matematico di Palermo (1884-1940)* 22.1 (1906), pp. 1–72.

[Pen99a] M. D. Penrose. "A strong law for the longest edge of the minimal spanning tree." In: *The Annals of Probability* 27.1 (1999), pp. 246–260.

[Pen99b] M. D. Penrose. "A strong law for the largest nearest-neighbour link between random points." In: *Journal of the london mathematical society* 60.3 (1999), pp. 951–960.

[TS15] N. G. Trillos and D. Slepčev. "On the rate of convergence of empirical measures in ∞-transportation distance." In: *Canadian Journal of Mathematics* 67.6 (2015), pp. 1358–1383.

[San15] F. Santambrogio. "Optimal transport for applied mathematicians." In: *Birkäuser, NY* 55.58-63 (2015), p. 94.

[BKB20] L. Bungert, Y. Korolev, and M. Burger. "Structural analysis of an L-infinity variational problem and relations to distance functions." In: *Pure and Applied Analysis* 2.3 (2020), pp. 703–738. DOI: 10.2140/paa.2020.2.703.

[Goo52] I. J. Good. "Rational decisions." In: *Journal of the Royal Statistical Society: Series B (Methodological)* 14.1 (1952), pp. 107–114.

[SGS15] R. K. Srivastava, K. Greff, and J. Schmidhuber. "Highway networks." In: *arXiv preprint arXiv:1505.00387* (2015).

[BREG-II] L. Bungert, T. Roith, D. Tenbrinck, and M. Burger. "Neural Architecture Search via Bregman Iterations." In: (2021). arXiv: 2106.02479 [cs.LG].

[Rie22] K. Riedl. "Leveraging memory effects and gradient information in consensus-based optimization: On global convergence in mean-field law." In: *arXiv preprint arXiv:2211.12184* (2022).

[Pin+17] R. Pinnau, C. Totzeck, O. Tse, and S. Martin. "A consensus-based model for global optimization and its mean-field limit." In: *Mathematical Models and Methods in Applied Sciences* 27.01 (2017), pp. 183–204.

[Car+21] J. A. Carrillo, S. Jin, L. Li, and Y. Zhu. "A consensus-based global optimization method for high dimensional machine learning problems." In: *ESAIM: Control, Optimisation and Calculus of Variations* 27 (2021), S5.

[Cau+47] A. Cauchy et al. "Méthode générale pour la résolution des systemes d'équations simultanées." In: *Comp. Rend. Sci. Paris* 25.1847 (1847), pp. 536–538.

[RM51] H. Robbins and S. Monro. "A stochastic approximation method." In: *The annals of mathematical statistics* (1951), pp. 400–407.

[Sha+18] A. Shafahi, W. R. Huang, C. Studer, S. Feizi, and T. Goldstein. "Are adversarial examples inevitable?" In: *arXiv preprint arXiv:1809.02104* (2018).

[FFF18] A. Fawzi, H. Fawzi, and O. Fawzi. "Adversarial vulnerability for any classifier." In: *Advances in neural information processing systems* 31 (2018).

[Sta+21] J. Stanczuk, C. Etmann, L. M. Kreusser, and C.-B. Schönlieb. "Wasserstein GANs work because they fail (to approximate the Wasserstein distance)." In: *arXiv preprint arXiv:2103.01678* (2021).

[Bun+23] L. Bungert, N. G. Trillos, M. Jacobs, D. McKenzie, Đ. Nikolić, and Q. Wang. "It begins with a boundary: A geometric view on probabilistically robust learning." In: *arXiv preprint arXiv:2305.18779* (2023).

[LC10] Y. LeCun and C. Cortes. "MNIST handwritten digit database." In: (2010).

[Eng+18] L. Engstrom, B. Tran, D. Tsipras, L. Schmidt, and A. Madry. "A rotation and a translation suffice: Fooling cnns with simple transformations." In: (2018).

[Guo+17] C. Guo, M. Rana, M. Cisse, and L. Van Der Maaten. "Countering adversarial images using input transformations." In: *arXiv preprint arXiv:1711.00117* (2017).

[ANR74] N. Ahmed, T. Natarajan, and K. R. Rao. "Discrete cosine transform." In: *IEEE transactions on Computers* 100.1 (1974), pp. 90–93.

[Yua+19] X. Yuan, P. He, Q. Zhu, and X. Li. "Adversarial examples: Attacks and defenses for deep learning." In: *IEEE transactions on neural networks and learning systems* 30.9 (2019), pp. 2805–2824.

[KGB16] A. Kurakin, I. Goodfellow, and S. Bengio. "Adversarial machine learning at scale." In: *arXiv preprint arXiv:1611.01236* (2016).

[Mad+17] A. Madry, A. Makelov, L. Schmidt, D. Tsipras, and A. Vladu. "Towards deep learning models resistant to adversarial attacks." In: *arXiv preprint arXiv:1706.06083* (2017).

[BGM23] L. Bungert, N. García Trillos, and R. Murray. "The geometry of adversarial training in binary classification." In: *Information and Inference: A Journal of the IMA* 12.2 (2023), pp. 921–968.

[LeC+95] Y. LeCun et al. "Learning algorithms for classification: A comparison on handwritten digit recognition." In: *Neural networks: the statistical mechanics perspective* 261.276 (1995), p. 2.

[Has+20] M. Hasannasab, J. Hertrich, S. Neumayer, G. Plonka, S. Setzer, and G. Steidl. "Parseval proximal neural networks." In: *Journal of Fourier Analysis and Applications* 26 (2020), pp. 1–31.

[Gou+20] H. Gouk, E. Frank, B. Pfahringer, and M. J. Cree. "Regularisation of neural networks by enforcing Lipschitz continuity." In: *Machine Learning* (2020), pp. 1–24.

[KMP20] V. Krishnan, A. A. A. Makdah, and F. Pasqualetti. "Lipschitz Bounds and Provably Robust Training by Laplacian Smoothing." In: *arXiv preprint arXiv:2006.03712* (2020).

[Sha+19] A. Shafahi, M. Najibi, M. A. Ghiasi, Z. Xu, J. Dickerson, C. Studer, L. S. Davis, G. Taylor, and T. Goldstein. "Adversarial training for free!" In: *Advances in Neural Information Processing Systems* 32 (2019).

[XRV17] H. Xiao, K. Rasul, and R. Vollgraf. "Fashion-MNIST: a Novel Image Dataset for Benchmarking Machine Learning Algorithms." In: *CoRR* abs/1708.07747 (2017).

[Hoe+21] T. Hoefler, D. Alistarh, T. Ben-Nun, N. Dryden, and A. Peste. "Sparsity in Deep Learning: Pruning and growth for efficient inference and training in neural networks." In: *J. Mach. Learn. Res.* 22.241 (2021), pp. 1–124.

[Gho+21] A. Gholami, S. Kim, Z. Dong, Z. Yao, M. W. Mahoney, and K. Keutzer. "A survey of quantization methods for efficient neural network inference." In: *arXiv preprint arXiv:2103.13630* (2021).

[EMH19] T. Elsken, J. H. Metzen, and F. Hutter. "Neural architecture search: A survey." In: *The Journal of Machine Learning Research* 20.1 (2019), pp. 1997–2017.

[How+17] A. G. Howard, M. Zhu, B. Chen, D. Kalenichenko, W. Wang, T. Weyand, M. Andreetto, and H. Adam. "Mobilenets: Efficient convolutional neural networks for mobile vision applications." In: *arXiv preprint arXiv:1704.04861* (2017).

[Ban+18] R. Banner, I. Hubara, E. Hoffer, and D. Soudry. "Scalable methods for 8-bit training of neural networks." In: *Advances in neural information processing systems* 31 (2018).

[CBD14] M. Courbariaux, Y. Bengio, and J.-P. David. "Training deep neural networks with low precision multiplications." In: *arXiv preprint arXiv:1412.7024* (2014).

[Sch92] J. Schmidhuber. "Learning complex, extended sequences using the principle of history compression." In: *Neural Computation* 4.2 (1992), pp. 234–242.

[HVD15] G. Hinton, O. Vinyals, and J. Dean. "Distilling the knowledge in a neural network." In: *arXiv preprint arXiv:1503.02531* (2015).

[LDS89] Y. LeCun, J. Denker, and S. Solla. "Optimal brain damage." In: *Advances in neural information processing systems* 2 (1989).

[CFP97] G. Castellano, A. M. Fanelli, and M. Pelillo. "An iterative pruning algorithm for feedforward neural networks." In: *IEEE transactions on Neural networks* 8.3 (1997), pp. 519–531.

[CM73] J. F. Claerbout and F. Muir. "Robust modeling with erratic data." In: *Geophysics* 38.5 (1973), pp. 826–844.

[Tib96] R. Tibshirani. "Regression shrinkage and selection via the lasso." In: *Journal of the Royal Statistical Society: Series B (Methodological)* 58.1 (1996), pp. 267–288.

[Nit14] A. Nitanda. "Stochastic proximal gradient descent with acceleration techniques." In: *Advances in Neural Information Processing Systems* 27 (2014), pp. 1574–1582.

[RVV20] L. Rosasco, S. Villa, and B. C. Vũ. "Convergence of stochastic proximal gradient algorithm." In: *Applied Mathematics & Optimization* 82.3 (2020), pp. 891–917.

[Moc+18] D. C. Mocanu, E. Mocanu, P. Stone, P. H. Nguyen, M. Gibescu, and A. Liotta. "Scalable training of artificial neural networks with adaptive sparse connectivity inspired by network science." In: *Nature communications* 9.1 (2018), pp. 1–12.

[DZ19] T. Dettmers and L. Zettlemoyer. "Sparse networks from scratch: Faster training without losing performance." In: *arXiv preprint arXiv:1907.04840* (2019).

[DYJ19] X. Dai, H. Yin, and N. K. Jha. "NeST: A neural network synthesis tool based on a grow-and-prune paradigm." In: *IEEE Transactions on Computers* 68.10 (2019), pp. 1487–1497.

[Fu+22] Y. Fu, C. Liu, D. Li, Z. Zhong, X. Sun, J. Zeng, and Y. Yao. "Exploring structural sparsity of deep networks via inverse scale spaces." In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* (2022).

[Liu+21] S. Liu, D. C. Mocanu, A. R. R. Matavalam, Y. Pei, and M. Pechenizkiy. "Sparse evolutionary deep learning with over one million artificial neurons on commodity hardware." In: *Neural Computing and Applications* 33.7 (2021), pp. 2589–2604.

[BB18] M. Benning and M. Burger. "Modern regularization methods for inverse problems." In: *Acta Numerica* 27 (2018), pp. 1–111.

[PB+14] N. Parikh, S. Boyd, et al. "Proximal algorithms." In: *Foundations and trends® in Optimization* 1.3 (2014), pp. 127–239.

[Sca+17] S. Scardapane, D. Comminiello, A. Hussain, and A. Uncini. "Group sparse regularization for deep neural networks." In: *Neurocomputing* 241 (2017), pp. 81–89.

[CP08] P. L. Combettes and J.-C. Pesquet. "Proximal thresholding algorithm for minimization over orthonormal bases." In: *SIAM Journal on Optimization* 18.4 (2008), pp. 1351–1376.

[DDD04] I. Daubechies, M. Defrise, and C. De Mol. "An iterative thresholding algorithm for linear inverse problems with a sparsity constraint." In: *Communications on Pure and Applied Mathematics: A Journal Issued by the Courant Institute of Mathematical Sciences* 57.11 (2004), pp. 1413–1457.

[FNW07] M. A. Figueiredo, R. D. Nowak, and S. J. Wright. "Gradient projection for sparse reconstruction: Application to compressed sensing and other inverse problems." In: *IEEE Journal of selected topics in signal processing* 1.4 (2007), pp. 586–597.

[Cha04] A. Chambolle. "An algorithm for total variation minimization and applications." In: *Journal of Mathematical imaging and vision* 20 (2004), pp. 89–97.

[CP11] A. Chambolle and T. Pock. "A first-order primal-dual algorithm for convex problems with applications to imaging." In: *Journal of mathematical imaging and vision* 40 (2011), pp. 120–145.

[Bre67] L. M. Bregman. "The relaxation method of finding the common point of convex sets and its application to the solution of problems in convex programming." In: *USSR computational mathematics and mathematical physics* 7.3 (1967), pp. 200–217.

[De 93] E. De Giorgi. "New problems on minimizing movements." In: *Ennio de Giorgi: Selected Papers* (1993), pp. 699–713.

[Osh+05] S. Osher, M. Burger, D. Goldfarb, J. Xu, and W. Yin. "An iterative regularization method for total variation-based image restoration." In: *Multiscale Modeling & Simulation* 4.2 (2005), pp. 460–489.

[ROF92] L. I. Rudin, S. Osher, and E. Fatemi. "Nonlinear total variation based noise removal algorithms." In: *Physica D: nonlinear phenomena* 60.1-4 (1992), pp. 259–268.

[Bur+06] M. Burger, G. Gilboa, S. Osher, J. Xu, et al. "Nonlinear inverse scale space methods." In: *Communications in Mathematical Sciences* 4.1 (2006), pp. 179–212.

[Bur+07] M. Burger, K. Frick, S. Osher, and O. Scherzer. "Inverse total variation flow." In: *Multiscale Modeling & Simulation* 6.2 (2007), pp. 366–395.

[Yin+08] W. Yin, S. Osher, D. Goldfarb, and J. Darbon. "Bregman iterative algorithms for ℓ_1-minimization with applications to compressed sensing." In: *SIAM Journal on Imaging sciences* 1.1 (2008), pp. 143–168.

[COS09] J.-F. Cai, S. Osher, and Z. Shen. "Convergence of the linearized Bregman iteration for ℓ_1-norm minimization." In: *Mathematics of Computation* 78.268 (2009), pp. 2127–2136.

[Vil+23] S. Villa, S. Matet, B. C. Vũ, and L. Rosasco. "Implicit regularization with strongly convex bias: Stability and acceleration." In: *Analysis and Applications* 21.01 (2023), pp. 165–191.

[BT03] A. Beck and M. Teboulle. "Mirror descent and nonlinear projected subgradient methods for convex optimization." In: *Operations Research Letters* 31.3 (2003), pp. 167–175.

[NY83] A. S. Nemirovskij and D. B. Yudin. "Problem complexity and method efficiency in optimization." In: (1983).

[Nem+09] A. Nemirovski, A. Juditsky, G. Lan, and A. Shapiro. "Robust stochastic approximation approach to stochastic programming." In: *SIAM Journal on optimization* 19.4 (2009), pp. 1574–1609.

[Nes83] Y. Nesterov. "A method for unconstrained convex minimization problem with the rate of convergence o (1/k2)." In: *Doklady ANSSSR* 269.3 (1983), pp. 543–547.

[Qia99] N. Qian. "On the momentum term in gradient descent learning algorithms." In: *Neural networks* 12.1 (1999), pp. 145–151.

[KB14] D. Kingma and J. Ba. "Adam: A Method for Stochastic Optimization." In: *arXiv preprint arXiv:1412.6980* (2014).

[HR21] F. Hanzely and P. Richtárik. "Fastest rates for stochastic mirror descent methods." In: *Computational Optimization and Applications* 79 (2021), pp. 717–766.

[ZH18] S. Zhang and N. He. "On the convergence rate of stochastic mirror descent for nonsmooth nonconvex optimization." In: *arXiv preprint arXiv:1806.04781* (2018).

[DOr+21] R. D'Orazio, N. Loizou, I. Laradji, and I. Mitliagkas. "Stochastic mirror descent: Convergence analysis and adaptive variants via the mirror stochastic Polyak stepsize." In: *arXiv preprint arXiv:2110.15412* (2021).

[AKL22] P.-C. Aubin-Frankowski, A. Korba, and F. Léger. "Mirror descent with relative smoothness in measure spaces, with application to sinkhorn and em." In: *Advances in Neural Information Processing Systems* 35 (2022), pp. 17263–17275.

[Ben+21] M. Benning, M. M. Betcke, M. J. Ehrhardt, and C.-B. Schönlieb. "Choose your path wisely: gradient descent in a Bregman distance framework." In: *SIAM Journal on Imaging Sciences* 14.2 (2021), pp. 814–843.

[HZ93] G. E. Hinton and R. Zemel. "Autoencoders, minimum description length and Helmholtz free energy." In: *Advances in neural information processing systems* 6 (1993).

[KLM21] N. Kovachki, S. Lanthaler, and S. Mishra. "On universal approximation and error bounds for Fourier neural operators." In: *The Journal of Machine Learning Research* 22.1 (2021), pp. 13237–13312.

[HW62] D. H. Hubel and T. N. Wiesel. "Receptive fields, binocular interaction and functional architecture in the cat's visual cortex." In: *The Journal of physiology* 160.1 (1962), p. 106.

[Li+20] Z. Li, N. Kovachki, K. Azizzadenesheli, B. Liu, K. Bhattacharya, A. Stuart, and A. Anandkumar. "Fourier neural operator for parametric partial differential equations." In: *arXiv preprint arXiv:2010.08895* (2020).

[Kov+21] N. B. Kovachki, Z. Li, B. Liu, K. Azizzadenesheli, K. Bhattacharya, A. M. Stuart, and A. Anandkumar. "Neural Operator: Learning Maps Between Function Spaces." In: *arXiv:2108.08481* (2021).

[Bri19] T. Briand. "Trigonometric Polynomial Interpolation of Images." In: *Image Processing On Line* 9 (Oct. 2019), pp. 291–316.

[HG16] D. Hendrycks and K. Gimpel. "Gaussian Error Linear Units (GELUs)." In: *arXiv:1606.08415* (2016).

Theses

[Roi21] T. Roith. "Master thesis: Continuum limit of Lipschitz learning on graphs." MA thesis. Friedrich-Alexander-Universität Erlangen-Nürnberg, 2021.

[Sma10] C. K. Smart. "On the infinity Laplacian and Hrushovski's fusion." PhD thesis. UC Berkeley, 2010.

[Kab22] S. Kabri. "Fourier Neural Operators for Image Classification." MA thesis. Friedrich-Alexander-Universität Erlangen-Nürnberg, 2022.

Online

[Pio21] G. Piosenka. *BIRDS 500 - SPECIES IMAGE CLASSIFICATION*. 2021. URL: %7Bhttps://www.kaggle.com/datasets/gpiosenka/100-bird-species%7D.

Part II.

Prints